Mathematics Bootcamp
Part II: Matrix Algebra and Multivariate Statistics

A. Zaidi[1]   K. Burris[1]

[1]Department of Statistical Science
Duke University

Graduate Orientation, August 2017

Handout and Slides

Please follow along with the handout made available in the github repository

https://github.com/DukeStatSci/MathBootcamp2017

Your team exercises will be available in the handout. I will make these slides (with the solutions) available after today's session.

Outline

Matrices: The Basics

Notation

A real **matrix** (hereafter referred to as a matrix) is a rectangular array of real numbers. The collection of real numbers within a matrix $a_{11}, a_{12}, \ldots, a_{1n}, \ldots, a_{m1}, a_{m2}, \ldots, a_{mn}$ can be arranged as

$$\begin{pmatrix} a_{11} & a_{12} & a_{13} & \ldots & a_{1n} \\ a_{21} & a_{22} & a_{23} & \ldots & a_{2n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & a_{m3} & \ldots & a_{mn} \end{pmatrix}$$

This matrix, which has m rows and n columns, is an $m \times n$ matrix, where m and n are the **dimensions** of the matrix. The scalar a_{ij} at the intersection of the ith row and jth column of a matrix is called the ijth entry or element. Boldface capital letters (e.g. **A**) are used to represent matrices.

Matrix Addition

The **sum** of two $m \times n$ matrices **A** and **B** is denoted by $\mathbf{A} + \mathbf{B}$ and defined to be the $m \times n$ matrix whose ijth element is $a_{ij} + b_{ij}$. For example,

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} + \begin{pmatrix} 5 & 6 \\ 7 & 8 \end{pmatrix} = \begin{pmatrix} 6 & 8 \\ 10 & 12 \end{pmatrix}$$

Addition is commutative and associative, so

$$\mathbf{A} + \mathbf{B} = \mathbf{B} + \mathbf{A}$$

$$\mathbf{A} + (\mathbf{B} + \mathbf{C}) = (\mathbf{A} + \mathbf{B}) + \mathbf{C}$$

Matrix Multiplication

Let **A** be a $m \times n$ matrix and let **B** be a $p \times q$ matrix. If $n = p$, the matrix **product AB** is defined to be the $m \times q$ matrix whose ijth element is

$$\mathbf{AB}_{ij} = \sum_{k=1}^{n} a_{ik} b_{kj}$$

For example,

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} \begin{pmatrix} 5 & 6 \\ 7 & 8 \end{pmatrix} = \begin{pmatrix} 1(5) + 2(7) & 1(6) + 2(8) \\ 3(5) + 4(7) & 3(6) + 4(8) \end{pmatrix} = \begin{pmatrix} 19 & 22 \\ 43 & 50 \end{pmatrix}$$

Matrix Multiplication Properties

Matrix multiplication is associative, but not commutative in general (**BA** may not even be defined!). That is, if a matrix **C** has q rows, then

$$\mathbf{A}(\mathbf{BC}) = (\mathbf{AB})\mathbf{C}$$

In addition, matrix multiplication is distributive with respect to addition, so assuming that matrix dimensions are such that all products and sums are defined,

$$\mathbf{A}(\mathbf{B} + \mathbf{C}) = \mathbf{AB} + \mathbf{AC}$$
$$(\mathbf{A} + \mathbf{B})\mathbf{C} = \mathbf{AC} + \mathbf{BC}$$

For any scalar c, we have that

$$c(\mathbf{A} + \mathbf{B}) = c\mathbf{A} + c\mathbf{B}$$
$$c\mathbf{AB} = (c\mathbf{A})\mathbf{B} = \mathbf{A}(c\mathbf{B})$$

Transposition

The transpose of a $m \times n$ matrix **A**, denoted by either the symbol **A**$^{'}$ or **A**^{T}, is the $n \times m$ matrix whose ijth element is the jith element of **A**. For example,

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix}^T = \begin{pmatrix} 1 & 3 \\ 2 & 4 \end{pmatrix}$$

Below are some properties of matrix transposes:

$$(\mathbf{A}^T)^T = \mathbf{A}$$

$$(\mathbf{A} + \mathbf{B})^T = \mathbf{A}^T + \mathbf{B}^T$$

$$(\mathbf{A}\mathbf{B})^T = \mathbf{B}^T\mathbf{A}^T$$

Square Matrices

A matrix that has the same number of rows as columns is called a **square** matrix. A $n \times n$ square matrix is said to be of order n. The n elements $a_{11}, a_{22}, \ldots, a_{nn}$ of

$$\mathbf{A} = \begin{pmatrix} a_{11} & a_{12} & \ldots & a_{1n} \\ a_{21} & a_{22} & \ldots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \ldots & a_{nn} \end{pmatrix}$$

that fall on the imaginary diagonal line extending from the top left corner to the bottom right corner of the matrix are known as **diagonal elements**. Any elements of the matrix that are not on the main diagonal are called **off-diagonal elements**. Note that the product **AA** is only defined if A is square. If it is, we can write $\mathbf{A}^2 = \mathbf{AA}$ and $\mathbf{A}^k = \mathbf{AAA} \cdots \mathbf{A}$.

Symmetric Matrices

A square matrix **A** is symmetric if $\mathbf{A}^T = \mathbf{A}$. For example, the matrix

$$\begin{pmatrix} 1 & 1 & 1 \\ 1 & 2 & 3 \\ 1 & 3 & 6 \end{pmatrix}$$

is symmetric.

Diagonal Matrices

A square matrix **A** is a **diagonal matrix** if all of its off-diagonal elements are equal to 0. That is,

$$\mathbf{A} = \begin{pmatrix} d_1 & 0 & \ldots & 0 \\ 0 & d_2 & \ldots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \ldots & d_n \end{pmatrix}$$

for some $d_1, d_2, \ldots, d_n$. Diagonal matrices are very easy to work with, since for a $m \times n$ matrix **A** and $n \times n$ diagonal matrix **D**, the ijth element of **DA** (or **AD**) is equal to $d_i a_{ij}$.

The Identity Matrix

The diagonal matrix

$$diag(1, 1, \ldots, 1) = \begin{pmatrix} 1 & 0 & \ldots & 0 \\ 0 & 1 & \ldots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \ldots & 1 \end{pmatrix}$$

with diagonal elements all equal to 1 is the identity matrix. We use the symbol $\mathbf{I}_n$ (or just $\mathbf{I}$ if the context is clear) to denote the identity matrix. For any matrix $\mathbf{A}$,

$$\mathbf{AI} = \mathbf{IA} = \mathbf{A}$$

Triangular Matrices

If all the elements of a square matrix that are located below and to the left of the diagonal are 0, then the matrix is called **upper triangular**. An upper triangular matrix looks like this:

$$\begin{pmatrix} a_{11} & a_{12} & a_{13} & \ldots & a_{1n} \\ 0 & a_{22} & a_{23} & \ldots & a_{2n} \\ 0 & 0 & a_{33} & \ldots & a_{3n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \ldots & a_{nn} \end{pmatrix}$$

Similarly, a matrix with all elements above and to the right of the main diagonal equal to zero is a **lower triangular matrix**. More formally, an $n \times n$ matrix **A** is upper triangular if $a_{ij} = 0$ for $j < i = 1, 2, \ldots, n$.

Trace of a Square Matrix

The **trace** of a square matrix **A** of order n is defined to be the sum of the n diagonal elements of **A**. This is denoted by the symbol $tr(\mathbf{A})$. Thus,

$$tr(\mathbf{A}) = a_{11} + a_{22} + \ldots + a_{nn}$$

Some properties of the trace are

$$tr(k\mathbf{A}) = k \; tr(\mathbf{A})$$
$$tr(\mathbf{A} + \mathbf{B}) = tr(\mathbf{A}) + tr(\mathbf{B})$$
$$tr(\mathbf{A}^T) = tr(\mathbf{A})$$
$$tr(\mathbf{A}\mathbf{B}) = tr(\mathbf{B}\mathbf{A})$$
$$tr(\mathbf{A}_1\mathbf{A}_2\cdots\mathbf{A}_k) = tr(\mathbf{A}_k\mathbf{A}_1\mathbf{A}_2\cdots\mathbf{A}_{k-1})$$

Vectors

A matrix with only one column, that is, a matrix of the form

$$\begin{pmatrix} a_1 \\ a_2 \\ \vdots \\ a_m \end{pmatrix}$$

is a **column vector**. Similarly, a matrix with one row is called a **row vector**. We distinguish vectors from matrices by referring to them by a boldface lowercase letter. So $\mathbf{a}$ is a column vector, and $\mathbf{a}^T$ is a row vector. The symbol $\mathbf{1}_m$ denotes a column vector with all entries equal to one.

Inner Product and Orthogonality

The **inner product** of two vectors $\mathbf{a}$ and $\mathbf{b}$ is $\mathbf{a}^T\mathbf{b}$. The sum of a vector $\mathbf{a}$ can be written as $\mathbf{1}^T\mathbf{a}$. The mean of a vector is

$$\mathbf{1}^T\mathbf{a}/n = (\mathbf{1}^T\mathbf{1})^{-1}\mathbf{1}^T\mathbf{a}$$

If $\mathbf{a}^T\mathbf{b} = 0$, then $\mathbf{a}$ and $\mathbf{b}$ are **perpendicular** to one another. A collection of vectors is **orthogonal** if and only if they are pairwise perpendicular.

A vector $\mathbf{a}$ is a **unit vector** if $\mathbf{a}^T\mathbf{a} = 1$.

Null vectors (and null matrices) are denoted by the notation $\mathbf{0}_m$.

Partitioned Matrices

A **submatrix** of a matrix **A** can be obtained by striking out rows and/or columns of **A**. A **partitioned matrix** is a $m \times n$ matrix **A** that has been expressed in the form

$$\mathbf{A} = \begin{pmatrix} \mathbf{A}_{11} & \mathbf{A}_{12} & \ldots & \mathbf{A}_{1c} \\ \mathbf{A}_{21} & \mathbf{A}_{22} & \ldots & \mathbf{A}_{2c} \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{A}_{r1} & \mathbf{A}_{r2} & \ldots & \mathbf{A}_{rc} \end{pmatrix}$$

where the submatrix $\mathbf{A}_{ij}$ is referred to as the *ij*th block of **A**. Partitioning matrices can be very useful, particularly when dealing with large sparse matrices (matrices that have many entries equal to zero).

Team Exercises for Unit 1

Problem 1

Observe that the *ij*th element of $(\mathbf{AB})^T$ equals the *ji*th element of $\mathbf{AB}$, which are both equal to

$$\sum_k a_{jk} b_{ki} = \sum_k b_{ki} a_{jk}$$

Since b_{ki} is the *ik*th element of $\mathbf{B}^T$ and a_{jk} is the *kj*th element of A^T, the *ij*th element of $\mathbf{B}^T \mathbf{A}^T$ is equal to the *ij*th element of $(\mathbf{AB})^T$.

Problem 2

By definition, the ijth element of **AB** is $\sum_{k=1}^{n} a_{ik} b_{kj}$. Suppose the **A** and **B** are both upper triangular. Then $a_{ik} = 0$ for $k < i$ and $b_{kj} = 0$ for $k > j$. Therefore,

$$\sum_{k=1}^{n} a_{ik} b_{kj} = \sum_{k=1}^{i-1} (0) b_{kj} + a_{ii} b_{ij} + \sum_{k=i+1}^{n} a_{ik}(0) = a_{ii} b_{ij}$$

Since **B** is upper triangular, $b_{ij} = 0$ for $j < i$ so **AB** is upper triangular. The diagonal elements of **AB** are $a_{ii} b_{ii}$.

Problem 3

Answers may vary. Let

$$\mathbf{A} = \begin{pmatrix} 1 & 2 \\ 2 & 3 \end{pmatrix} \qquad \mathbf{B} = \begin{pmatrix} 1 & 2 \\ 2 & 1 \end{pmatrix}$$

Then

$$\mathbf{AB} = \begin{pmatrix} 5 & 8 \\ 4 & 7 \end{pmatrix}$$

which is not symmetric. In general, the product of two symmetric matrices will be symmetric if and only if the two matrices commute. To see this, note that if **AB** is symmetric,

$$\mathbf{AB} = (\mathbf{AB})^T = \mathbf{B}^T \mathbf{A}^T = \mathbf{BA}$$

If **B** and **A** commute, then

$$\mathbf{AB} = \mathbf{BA} = \mathbf{B}^T \mathbf{A}^T = (\mathbf{AB})^T$$

Matrices: Beyond Basic Operations

Linear Independence

A nonempty finite set $\{\mathbf{A}_1, \mathbf{A}_2, \ldots, \mathbf{A_k}\}$ of $m \times n$ matrices is **linearly dependent** if there exist scalars $x_1, x_2, \ldots, x_k$, not all equal to zero, such that

$$x_1 \mathbf{A}_1 + x_2 \mathbf{A_2} + \cdots + x_k \mathbf{A}_k = \mathbf{0}$$

If no such scalars exist, the set is said to be **linearly independent**.

In addition, a set of two or more $m \times n$ matrices is linearly dependent if and only if at least one of the matrices can be expressed as a linear combination of the other matrices ($\mathbf{A_j}$ is expressible as a linear combination of $\mathbf{A_1}, \ldots, \mathbf{A}_{j-1}, \mathbf{A}_{j+1}, \ldots \mathbf{A_k}$).

Row and Column Space

The **column space** of a $m \times n$ **A** is the set whose elements consist of all m dimensional vectors that can be expressed as linear combinations of the n columns of **A**. The elements of the column space of **A** are all m dimensional vectors of the form

$$x_1\mathbf{a}_1 + x_2\mathbf{a}_2 + \cdots + x_n\mathbf{a}_n$$

where $x_1, x_2, \ldots, x_n$ are scalars and $\mathbf{a}_1, \ldots, \mathbf{a}_n$ are the columns of **A**. The row space is defined similarly, by finding all n dimensional vectors that can be expressed as a linear combination of the rows of **A**.

Column Space: Example

For example, the column space of the 3×4 matrix

$$\begin{pmatrix} 2 & -4 & 0 & 0 \\ -1 & 2 & 0 & 0 \\ 0 & 0 & 1 & 2 \end{pmatrix}$$

includes the column vector

$$\begin{pmatrix} 4 \\ -2 \\ -3 \end{pmatrix} = 2 \begin{pmatrix} 2 \\ -1 \\ 0 \end{pmatrix} + 0 \begin{pmatrix} -4 \\ 2 \\ 0 \end{pmatrix} - 3 \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix} + 0 \begin{pmatrix} 0 \\ 0 \\ 2 \end{pmatrix}$$

but not the column vector $(2, 0, 0)^T$.

Linear Spaces

Row spaces and column spaces of a matrix are examples of a more general concept called a **linear space**. A nonempty set $\mathcal{V}$ is a linear space if it is closed under element-wise multiplication and scalar multiplication (i.e. $\mathbf{A}, \mathbf{B} \in \mathcal{V} \Rightarrow \mathbf{A} + \mathbf{B} \in \mathcal{V}, k\mathbf{A} \in \mathcal{V}$).

A subset $\mathcal{U}$ of a linear space $\mathcal{V}$ is called a **subspace** of $\mathcal{V}$ if $\mathcal{U}$ is a linear space. For example, the column space $\mathcal{C}(\mathbf{A})$ of a $m \times n$ matrix $\mathbf{A}$ is a subspace of $\mathbb{R}^m$ and the row space $\mathcal{C}(\mathbf{A})$ of $\mathbf{A}$ is a subspace of $\mathbb{R}^n$.

Basis

The **span** of a finite, nonempty set of matrices
$S = \{\mathbf{A}_1, \mathbf{A}_2, \ldots, \mathbf{A}_k\}$ is the set of all matrices that are expressible
as a linear combination of the elements of S. The span is denoted
by $\text{sp}(S)$ and is a linear space.

A **basis** for a linear space $\mathcal{V}$ is a finite set of linearly independent
matrices in $\mathcal{V}$ that spans $\mathcal{V}$. Note that if the columns of $\mathbf{A}$ are
linearly independent, then the columns of $\mathbf{A}$ are a basis for $\mathcal{C}(\mathbf{A})$.

For example, $\{(1,0,0)^T, (0,1,0)^T, (0,0,1)^T\}$ and
$\{(1,1,0)^T, (1,2,0)^T, (0,2,3)^T\}$ are both bases for $\mathbb{R}^3$. However,
$\{(1,0,0)^T, (0,1,0)^T, (0,0,1)^T, (0,2,3)^T\}$ is not a basis for $\mathbb{R}^3$
since the set is linearly dependent.

Rank

The number of matrices in a basis for a linear space $\mathcal{V}$ is called the **dimension** of $\mathcal{V}$ and is denoted by $\dim(\mathcal{V})$. The **row rank** of a matrix $\mathbf{A}$ is the dimension of the row space of $\mathbf{A}$. Similarly, the **column rank** of a matrix $\mathbf{A}$ is the dimension of the column space of $\mathbf{A}$.

An important result in matrix algebra is the row rank and the column of a matrix are equal. This common value is called the **rank** of $\mathbf{A}$ and denoted by $\text{rank}(\mathbf{A})$.

A matrix is said to have **full row rank** if $\text{rank}(\mathbf{A}) = m$ and **full column rank** if $\text{rank}(\mathbf{A}) = n$. Square matrices ($m = n$) are said to be **full rank** or **nonsingular** if it has both full row rank and full column rank. A $n \times n$ matrix with rank less than n is **singular**.

Rank: Example

$$\mathbf{A} = \begin{pmatrix} 1 & 3 & 4 \\ 2 & 6 & 8 \\ 5 & 7 & 12 \end{pmatrix}$$

is a singular matrix, since rank($\mathbf{A}$) = 2, but $\mathbf{A}$ is a 3×3 matrix.

Matrix Inverses

A $n \times n$ matrix $\mathbf{A}$ is invertible if and only if $\mathbf{A}$ is nonsingular. If $\mathbf{A}$ is invertible, then there exists a unique inverse matrix $\mathbf{A}^{-1}$ such that $\mathbf{A}\mathbf{A}^{-1} = \mathbf{A}^{-1}\mathbf{A} = \mathbf{I}$.

For any $n \times n$ nonsingular diagonal matrix $\mathbf{D}$,

$$\mathbf{D}^{-1} = diag(1/d_1, 1/d_2, \ldots, 1/d_n)$$

The inverse of a 2×2 nonsingular matrix

$$\mathbf{A} = \begin{pmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{pmatrix}$$

$$\mathbf{A}^{-1} = (1/k) \begin{pmatrix} a_{22} & -a_{12} \\ -a_{21} & a_{11} \end{pmatrix}$$

where $k = a_{22}a_{11} - a_{12}a_{21}$.

Inverse Properties

Let **A** be a nonsingular $n \times n$ matrix. Then

$$(\mathbf{A}^{-1})^{-1} = \mathbf{A}$$

$$(k\mathbf{A})^{-1} = 1/k\mathbf{A}^{-1}$$

for any non-zero scalar k. If **B** is also a nonsingular $n \times n$ matrix, then

$$((\mathbf{A}\mathbf{B})^{-1} = \mathbf{B}^{-1}\mathbf{A}^{-1}$$

and more generally,

$$(\mathbf{A_1}\mathbf{A_2}\cdots\mathbf{A_k})^{-1} = \mathbf{A_k}^{-1}\cdots\mathbf{A_2}^{-1}\mathbf{A_1}^{-1}$$

Orthogonal Matrices

A nonsingular matrix $\mathbf{A}$ is an **orthogonal matrix** if $\mathbf{A}^T = \mathbf{A}^{-1}$.
Equivalently, $\mathbf{A}^T\mathbf{A} = \mathbf{A}\mathbf{A}^T = \mathbf{I}$

Examples of orthogonal matrices include the identity matrix $\mathbf{I}_n$
and, for any angle θ, the 2×2 matrix

$$\begin{pmatrix} \cos(\theta) & \sin(\theta) \\ -\sin(\theta) & \cos(\theta) \end{pmatrix}$$

Generalized Inverses

A **generalized inverse** or **pseudoinverse** of an $m \times n$ matrix **A** is any $n \times m$ matrix **G** such that

$$\mathbf{AGA} = \mathbf{A}$$

.

If **A** is nonsingular, then it has a unique generalized inverse. Otherwise, **A** has an infinite number of generalized inverses. There are many properties of generalized inverses that are useful when solving a linear system, but they are not presented here for brevity.

Determinant

The **determinant** of a square $n \times n$ matrix $\mathbf{A}$ is denoted by either $|\mathbf{A}|$ or $\det(\mathbf{A})$. The minor $\mathbf{M}_{ij}$ of element A_{ij} is the $n-1 \times n-1$ matrix that is formed by removing the ith row and jth column from $\mathbf{A}$. The cofactor of A_{ij} is $C_{ij} = (-1)^{i+j}|\mathbf{M}_{ij}|$. Expanding along the ith row,

$$|\mathbf{A}| = \sum_{j=1}^{n} A_{ij} C_{ij}$$

Note that $|\mathbf{A}^T| = |\mathbf{A}|$ and that $|k\mathbf{A}| = k^n|\mathbf{A}|$. $\mathbf{A}$ is singular if $|\mathbf{A}| = 0$ and nonsingular otherwise.

Projection Matrices

A square matrix **A** is **idempotent** if $\mathbf{A}^2 = \mathbf{A}$. A square matrix **P** is a **projection matrix** if and only if **P** is idempotent. If **P** is also orthogonal, then **P** is called an orthogonal projector.

A very useful property of projection matrices is that their rank is equal to their trace, namely

$$tr(\mathbf{A}) = rank(\mathbf{A})$$

The identity matrix $\mathbf{I}_n$ is the only full rank idempotent matrix.

Application to Least Squares

A primary application of projection matrices is linear models and the method of least squares. One important result is that the projection $\mathbf{z}$ of a n-dimensional vector $\mathbf{y}$ onto the column space of a $n \times p$ matrix $\mathbf{X}$ is

$$\mathbf{z} = \mathbf{X}(\mathbf{X^T X})^{-1}\mathbf{X^T y}$$

Letting $\hat{\beta}$ be the ordinary least squares estimate of β, where $\mathbf{y} = \mathbf{X}\beta + \epsilon$, we have that the fitted values $\mathbf{X}\hat{\beta}$ is the projection of the response vector $\mathbf{y}$ onto the column space of the covariate matrix $\mathbf{X}$.

Quadratic Forms

If $\mathbf{A}$ is an $n \times n$ matrix and $\mathbf{x}$ is an n-dimensional vector, then a **quadratic form** is

$$\mathbf{x}^T \mathbf{A} \mathbf{x}$$

$\mathbf{A}$ is **positive definite** if, for all $\mathbf{x} \neq \mathbf{0}$, $\mathbf{x}^T \mathbf{A} \mathbf{x} > 0$. If instead $\mathbf{x}^T \mathbf{A} \mathbf{x} \geq 0$, then $\mathbf{A}$ is **nonnegative definite**.

Symmetric nonnegative definite matrices are encountered frequently in linear models and other areas of statistics. In particular, the covariance matrix (we'll get to that next section) is nonnegative definite. In addition, any positive definite matrix is nonsingular.

Eigenvalues and Eigenvectors

A scalar λ is said to be an **eigenvalue** of an $n \times n$ matrix **A** if there exists a non-null vector **x** such that

$$\mathbf{A}\mathbf{x} = \lambda\mathbf{x}$$

The set of eigenvalues is called the **spectrum** of **A**. All eigenvalues of positive-definite matrices are positive and all eigenvalues of nonnegative-definite matrices are nonnegative.

A non-null vector **x** is an **eigenvector** of **A** if there exists a scalar λ such that $\mathbf{A}\mathbf{x} = \lambda\mathbf{x}$.

Eigendecomposition

If **A** is a symmetric $n \times n$ matrix, eigenvectors $\mathbf{v}_j$ and $\mathbf{v}_k$ associated with distinct eigenvalues $\lambda_j \neq \lambda_k$ are orthogonal. In addition, if $\mathbf{V} = (\mathbf{v}_1, \mathbf{v}_2, \ldots, \mathbf{v_n})$, then the spectral theorem expresses **A** as a weighted average of rank 1 matrices,

$$\mathbf{A} = \mathbf{V}\mathbf{\Lambda}\mathbf{V}^T = \sum_{j=1}^{n} \lambda_j \mathbf{v}_j \mathbf{v}_j^T$$

This decomposition of a square matrix into its eigenvalues and eigenvectors is called the **eigendecomposition**. This decomposition will always exist if the matrix is symmetric. The rank of **A** is the number of nonzero eigenvalues, and

$$tr(\mathbf{A}) = \sum_{j=1}^{n} \lambda_j \qquad |\mathbf{A}| = \prod_{j=1}^{n} \lambda_j$$

Matrix Decompositions

The Cholesky decomposition takes a positive definite matrix $\mathbf{A}$ and decomposes it into $\mathbf{U}^T\mathbf{U}$, where $\mathbf{U}$ is an upper triangular matrix. Cholesky factorization is substantially faster than the more general LU factorization because pivoting is not required.

The QR decomposition decomposes a general $m \times n$ matrix $\mathbf{A}$ into $\mathbf{QR}$, where $\mathbf{Q}$ is an orthogonal $m \times m$ matrix and $\mathbf{R}$ is an upper triangular $m \times n$ matrix. This decomposition is useful for solving least squares problems and numerically approximating eigenvectors.

Singular Value Decomposition

A $m \times n$ matrix **A** can be expressed as

$$\mathbf{A} = \mathbf{U}\mathbf{D}\mathbf{V}^T$$

where **U** is an $m \times n$ matrix such that $\mathbf{U}^T\mathbf{U} = \mathbf{I_m}$, $\mathbf{V}^T\mathbf{V} = \mathbf{I_n}$, and **D** is a nonnegative diagonal matrix with entries $d_1 \geq \cdots \geq d_n$.

This representation is called the **singular value decomposition** of **A** (or SVD). The SVD provides a description of the within-row variation via **V** and a description of the within column variation via **U**. This is very useful when describing large data matrices with a low-dimensional approximation.

The SVD of a symmetric matrix is identical to its eigendecomposition.

The Jacobian Matrix

Let $\mathbf{x} = (x_1, \ldots, x_K)^T$ be a k dimensional vector and let

$$\mathbf{y} = (y_1, \ldots, y_J)^T = (f_1(\mathbf{x}), \ldots, f_J(\mathbf{x}))^T = \mathbf{f}(\mathbf{x})$$

be a J dimensional vector, where $\mathbf{f} : \mathbb{R}^K \to \mathbb{R}^J$. Then the partial derivative of $\mathbf{y}$ with respect to $\mathbf{x}^T$ yields the $J \times K$ **Jacobian matrix**

$$\mathbf{J_x y} = \frac{\partial \mathbf{y}}{\partial \mathbf{x}^T} = \begin{pmatrix} \frac{\partial y_1}{\partial x_1} & \frac{\partial y_1}{\partial x_2} & \cdots & \frac{\partial y_1}{\partial x_K} \\ \frac{\partial y_2}{\partial x_1} & \frac{\partial y_2}{\partial x_2} & \cdots & \frac{\partial y_2}{\partial x_K} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial y_J}{\partial x_1} & \frac{\partial y_J}{\partial x_2} & \cdots & \frac{\partial y_J}{\partial x_K} \end{pmatrix}$$

The **Jacobian** of the tranformation $\mathbf{y} = \mathbf{f}(\mathbf{x})$ is

$$J = |\mathbf{J_x y}|$$

Gradient and Hessian

If $y = f(\mathbf{x})$ is a scalar, then the **gradient** vector is

$$\Delta_x y = \frac{\partial y}{\partial \mathbf{x}} = \left(\frac{\partial y}{\partial x_1}, \frac{\partial y}{\partial x_2}, \ldots, \frac{\partial y}{\partial x_K} \right)^T = \left(\frac{\partial y}{\partial \mathbf{x}^T} \right)^T = (\mathbf{J}_x y)^T$$

The $K \times K$ matrix of second order partial derivatives is called the **Hessian matrix**. This is defined mathematically as

$$\mathbf{H}_x y = \frac{\partial}{\partial \mathbf{x}} \left(\frac{\partial y}{\partial \mathbf{x}} \right)^T = \frac{\partial^2 y}{\partial \mathbf{x} \partial \mathbf{x}^T} = \begin{pmatrix} \frac{\partial^2 y}{\partial x_1^2} & \frac{\partial^2 y}{\partial x_1 x_2} & \cdots & \frac{\partial^2 y}{\partial x_1 x_K} \\ \frac{\partial^2 y}{\partial x_1 x_2} & \frac{\partial^2 y}{\partial x_2^2} & \cdots & \frac{\partial^2 y}{\partial x_2 x_K} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 y}{\partial x_1 x_K} & \frac{\partial^2 y}{\partial x_2 x_K} & \cdots & \frac{\partial^2 y}{\partial x_K^2} \end{pmatrix}$$

The gradient and Hessian, if they can be found in closed form, are very useful when optimizing a scalar function $y = f(\mathbf{x})$.

Matrix Differentiation

The derivative of a $J \times K$ matrix **A** with respect to an r dimensional vector x is the $(Jr \times K)$ matrix of derivatives of **A** with respect to each element of **x**. In other words,

$$\frac{\partial \mathbf{A}}{\partial \mathbf{x}} = \left(\frac{\partial \mathbf{A}^T}{\partial x_1}, \ldots, \frac{\partial \mathbf{A}^T}{\partial x_r} \right)^T$$

Similar to normal differentiation, we have the following rules:

$$\frac{\partial(\alpha \mathbf{A})}{\partial \mathbf{x}} = \alpha \frac{\partial \mathbf{A}}{\partial \mathbf{x}}$$

$$\frac{\partial(\mathbf{A} + \mathbf{B})}{\partial \mathbf{x}} = \frac{\partial \mathbf{A}}{\partial \mathbf{x}} + \frac{\partial \mathbf{B}}{\partial \mathbf{x}}$$

$$\frac{\partial(\mathbf{A}\mathbf{B})}{\partial \mathbf{x}} = \left(\frac{\partial \mathbf{A}}{\partial \mathbf{x}} \right) \mathbf{B} + \mathbf{A} \left(\frac{\partial \mathbf{B}}{\partial \mathbf{x}} \right)$$

Team Exercises for Unit 2

Problem 1

Let **P** and **Q** both be $n \times n$ orthogonal matrices. Then

$$(\mathbf{P}\mathbf{Q})^{-1} = \mathbf{Q}^{-1}\mathbf{P}^{-1} = \mathbf{Q}^{T}\mathbf{P}^{T} = (\mathbf{P}\mathbf{Q})^{T}$$

Problem 2

If a matrix is idempotent, then it is a projection matrix.

$$(\mathbf{I} - \mathbf{P})(\mathbf{I} - \mathbf{P}) = \mathbf{I}^2 - 2\mathbf{I}\mathbf{P} + \mathbf{P}^2 = \mathbf{I} - 2\mathbf{P} + \mathbf{P} = \mathbf{I} - \mathbf{P}$$

Hence, $\mathbf{I} - \mathbf{P}$ is idempotent and it is therefore a projection matrix.

Problem 3

Let λ be an eigenvalue of an idempotent matrix $\mathbf{A}$. Then there exists an eigenvector $\mathbf{x}$ such that $\mathbf{A}\mathbf{x} = \lambda\mathbf{x}$. In addition, since $\mathbf{A} = \mathbf{A}^2$, $\lambda^2\mathbf{x} = \lambda\mathbf{x}$. Therefore, since $\lambda = \lambda^2$, λ must be either 0 or 1.

Problem 4

The centering matrix is

$$\mathbf{C} = (\mathbf{I} - \mathbf{1}(\mathbf{1}^T\mathbf{1})^{-1}\mathbf{1}^T)$$

It is not difficult to show that $\mathbf{P} = \mathbf{1}(\mathbf{1}^T\mathbf{1})^{-1}\mathbf{1}^T$ is a projection matrix and from exercise 2, we know that $\mathbf{I} - \mathbf{P}$ is also a projection matrix.

Problem 5

The Jacobian matrix in this case is

$$\begin{pmatrix} \frac{\partial x}{\partial r} & \frac{\partial x}{\partial \theta} \\ \frac{\partial y}{\partial r} & \frac{\partial y}{\partial \theta} \end{pmatrix} = \begin{pmatrix} \cos \theta & -r \sin \theta \\ \sin \theta & r \cos \theta \end{pmatrix}$$

Random Matrices and Multivariate Statistics

Random Vectors

If we have d random variables $X_1, X_2, \ldots, X_d$, each defined on the real line, we can write them as the d dimensional column vector

$$\mathbf{X} = (X_1, \cdots X_d)^T$$

which we call a d-dimensional **random vector**. The joint distribution function of the random vector $\mathbf{X}$ is

$$
\begin{aligned}
F_X(\mathbf{x}) &= F_X(x_1, \ldots, x_d) \\
&= P(X_1 \leq x_1, \ldots, X_d \leq x_d) \\
&= P(\mathbf{X} \leq \mathbf{x})
\end{aligned}
$$

If F_X is absolutely continuous, then the joint density function f_X of $\mathbf{X}$ is

$$f_X(\mathbf{x}) = f_X(x_1, \ldots, x_d) = \frac{\partial^d F_X(x_1, \ldots, x_d)}{\partial x_1 \cdots \partial x_d}$$

Random Vectors

To find the marginal density of a subset of the d variables, you can just integrate the others out. For example, if we have a joint bivariate density $f_{X_1,X_2}(x_1, x_2)$, then

$$f_{X_1}(x_1) = \int_{-\infty}^{\infty} f_{X_1,X_2}(x_1, x_2)dx_2 \qquad f_{X_2}(x_2) = \int_{-\infty}^{\infty} f_{X_1,X_2}(x_1, x_2)dx_1$$

The components of a random vector **X** are **independent** if the joint distribution function is a product of the marginal distribution functions

$$F_X(\mathbf{x}) = \prod_{i=1}^{d} F_i(x_i)$$

In addition, the joint density is the product of marginals

$$f_X(\mathbf{x}) = \prod_{i=1}^{d} f_i(x_i)$$

Expectation and Covariance

If **X** is a random vector with values in $\mathbb{R}^d$, then its expected value is given by the d dimensional vector

$$\mu_X = E(\mathbf{X}) = (E(X_1), \cdots, E(X_d)) = (\mu_1, \cdots, \mu_d)^T$$

and the $d \times d$ **covariance matrix** of **X** is

$$
\begin{aligned}
\boldsymbol{\Sigma}_{XX} &= \operatorname{cov}(\mathbf{X}, \mathbf{X}) \\
&= E[(\mathbf{X} - \mu_X)(\mathbf{X} - \mu_X)^T] \\
&= E[(X_1 - \mu_1, \cdots, X_d - \mu_d)(X_1 - \mu_1, \cdots, X_d - \mu_d)^T] \\
&= \begin{pmatrix} \sigma_1^2 & \sigma_{12} & \ldots & \sigma_{1d} \\ \sigma_{21} & \sigma_2^2 & \ldots & \sigma_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ \sigma_{d1} & \sigma_{d2} & \ldots & \sigma_d^2 \end{pmatrix}
\end{aligned}
$$

Correlation Matrix

The **correlation matrix** of **X** can be obtained by from $\boldsymbol{\Sigma}_{XX}$ by dividing the ith row by σ_i and the jth column by σ_j. The $d \times d$ matrix is then

$$P_{XX} = \begin{pmatrix} 1 & \rho_{12} & \ldots & \rho_{1d} \\ \rho_{21} & 1 & \ldots & \rho_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ \rho_{d1} & \rho_{d2} & \ldots & 1 \end{pmatrix}$$

where

$$\rho_{ij} = \rho_{ji} = \begin{cases} \frac{\sigma_{ij}}{\sigma_i \sigma_j} & i \neq j \\ 1 & \text{otherwise} \end{cases}$$

is the pairwise correlation coefficient between X_i and X_j. The correlation coefficient will always lie between -1 and 1 and is a measure of association between X_i and X_j.

Linear Functions of Random Vectors

If **Y** is a linear function of **X** such that

$$\mathbf{Y} = \mathbf{A}\mathbf{X} + \mathbf{b}$$

the mean vector and covariance matrix of **Y** is given by

$$\mu_Y = \mathbf{A}\mu_x + \mathbf{b}$$
$$\mathbf{\Sigma}_{YY} = \mathbf{A}\mathbf{\Sigma}_{XX}\mathbf{A}^T$$

Multivariate Normal Distribution

The form of the multivariate normal looks similar to that of the univariate normal. A random d vector $\mathbf{X}$ follows a multivariate normal distribution with mean vector μ and positive definite symmetric covariance matrix $\boldsymbol{\Sigma}$ if it has the density function

$$f(\mathbf{x}|\mu, \boldsymbol{\Sigma}) = (2\pi)^{-d/2} |\boldsymbol{\Sigma}|^{-1/2} e^{-\frac{1}{2}(\mathbf{x}-\mu)^T \boldsymbol{\Sigma}^{-1}(\mathbf{x}-\mu)}$$

We notationally denote a d dimensional normal distribution as

$$\mathbf{X} \sim N_d(\mu, \boldsymbol{\Sigma})$$

Multivariate Normal Distribution

The **Mahalanobis distance** from $\mathbf{x}$ to μ is given by the quadratic form

$$\Delta = \sqrt{(\mathbf{x} - \mu)^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \mu)}$$

An important result is that a random vector $\mathbf{X}$ follows a multivariate distribution if and only if every linear function of $\mathbf{X}$ follows a univariate normal distribution.

In linear models, we often assume that $\boldsymbol{\Sigma} = \sigma^2 \mathbf{I_d}$, in which case the density function reduces to

$$f(\mathbf{x}|\mu, \sigma) = (2\pi\sigma)^{-d/2} e^{-\frac{1}{2}(\mathbf{x} - \mu)^T (\mathbf{x} - \mu)}$$

Partitioned Random Vectors

Suppose we have two random vectors $\mathbf{X}$ and $\mathbf{Y}$, where $\mathbf{X}$ has d_1 components and $\mathbf{Y}$ had d_2 components. Let $\mathbf{Z}$ be the random $d_1 + d_2$ vector

$$\mathbf{Z} = \begin{pmatrix} \mathbf{X} \\ \mathbf{Y} \end{pmatrix}$$

Then the expected value and covariance matrix of $\mathbf{Z}$ is given by

$$\mu_Z = E[\mathbf{Z}] = \begin{pmatrix} E[\mathbf{X}] \\ E[\mathbf{Y}] \end{pmatrix} = \begin{pmatrix} \mu_X \\ \mu_Y \end{pmatrix}$$

$$\mathbf{\Sigma}_{ZZ} = \begin{pmatrix} cov(\mathbf{X}, \mathbf{X}) & cov(\mathbf{X}, \mathbf{Y}) \\ cov(\mathbf{Y}, \mathbf{X}) & cov(\mathbf{Y}, \mathbf{Y}) \end{pmatrix}$$

$$= \begin{pmatrix} \mathbf{\Sigma_{XX}} & \mathbf{\Sigma}_{XY} \\ \mathbf{\Sigma}_{YX} & \mathbf{\Sigma}_{YY} \end{pmatrix}$$

where $\mathbf{\Sigma}_{XY} = \mathbf{\Sigma}_{YX}^T$.

Marginal/Conditional Normal Distribution

The marginal distribution of **Y** is

$$\mathbf{Y} \sim N_{d_2}(\mu_y, \mathbf{\Sigma}_{YY}$$

The conditional distribution of **Y** given that **X** = **x** is multivariate normal with mean vector and covariance matrix given by

$$\mu_{Y|X} = \mu_Y + \mathbf{\Sigma}_{YX}\mathbf{\Sigma}_{XX}^{-1}(\mathbf{x} - \mu_X)$$
$$\mathbf{\Sigma}_{Y|X} = \mathbf{\Sigma}_{YY} - \mathbf{\Sigma}_{YX}\mathbf{\Sigma}_{XX}^{-1}\mathbf{\Sigma}_{XY}$$

Wishart Distribution

Given n independent and identically distributed d vectors

$$\mathbf{X}_i \sim N_d(\mu, \mathbf{\Sigma})$$

we say that the random positive-definite, symmetric matrix

$$\mathbf{W} = \sum_{i=1}^{n} \mathbf{X_i X_i}^T$$

follows a **Wishart distribution** with n degrees of freedom and matrix $\mathbf{\Sigma}$. We denote the Wishart distribution by

$$\mathbf{W} \sim \mathcal{W}_d(n, \mathbf{\Sigma})$$

You can think of the Wishart as a randomly drawn covariance matrix multiplied by the degrees of freedom n, since $E[\mathbf{W}] = n\mathbf{\Sigma}$. As $n \to \infty$, $\mathbf{W}/n \to \mathbf{\Sigma}$.

Properties of the Wishart Distribution

1. Let $\mathbf{W}_j \sim \mathcal{W}_d(n_j, \boldsymbol{\Sigma})$ be independent. Then $\sum_{j=1}^m \mathbf{W}_j \sim \mathcal{W}_d(\sum_{j=1}^m n_j, \boldsymbol{\Sigma})$

2. Suppose $\mathbf{W} \sim \mathcal{W}_d(n, \boldsymbol{\Sigma})$ and let $\mathbf{A}$ be a constant matrix having full row rank. Then $\mathbf{A}\mathbf{W}\mathbf{A}^T \sim \mathcal{W}_d(n, \mathbf{A}\boldsymbol{\Sigma}\mathbf{A}^T)$.

3. Suppose $\mathbf{W} \sim \mathcal{W}_d(n, \boldsymbol{\Sigma})$ and let $\mathbf{a}$ be a fixed d dimensional vector. Then $\mathbf{a}^T\mathbf{W}\mathbf{a} \sim (\mathbf{a}^T\boldsymbol{\Sigma}\mathbf{a})\chi_n^2$.

You can think of the Wishart as a multidimensional chi-square distribution. If $\mathbf{W}$ follows a Wishart distribution, then $\mathbf{W}^{-1}$ follows an **inverse Wishart distribution**.

Team Exercises for Unit 3

Solutions

1. Find the correlation matrix ρ of **X**

$$\rho = \begin{pmatrix} 1 & 0 & 1/6 \\ 0 & 1 & 0 \\ 1/6 & 0 & 1 \end{pmatrix}$$

2. Find the marginal distribution of X_2.

$$X_2 \sim N(1,1)$$

3. Find the marginal distribution of $\{X_1, X_3\}$.

$$\{X_1, X_3\} \sim N\left(\begin{pmatrix} 0 \\ -1 \end{pmatrix}, \begin{pmatrix} 3 & 1 \\ 1 & 2 \end{pmatrix} \right)$$

Solutions

4. Find the conditional distribution of $X_1 | X_3 = -1$.

   Using the conditional distribution formula, the conditional distribution of $\{X_1, X_2\}$ given $X_3 = -1$ is

   $$\mu = \begin{pmatrix} 0 \\ 1 \end{pmatrix} + \begin{pmatrix} 1 \\ 0 \end{pmatrix} (1/2)(0) = \begin{pmatrix} 0 \\ 1 \end{pmatrix}$$
   $$\boldsymbol{\Sigma} = \begin{pmatrix} 3 & 0 \\ 0 & 1 \end{pmatrix} - \begin{pmatrix} 1 \\ 0 \end{pmatrix} (1/2) \begin{pmatrix} 1 \\ 0 \end{pmatrix}^T = \begin{pmatrix} 5/2 & 0 \\ 0 & 1 \end{pmatrix}$$

   So looking at the marginal, the conditional distribution of X_1 is $N(0, 5/2)$.

5. Find the conditional distribution of $X_1 | \{X_2 = 1, X_3 = -1\}$
   $\{X_1, X_2\}$ given $X_3 = -1$ is
   You can do this using the conditional distribution formula or note that X_1 is independent of X_2 (from the next question). So the answer will be the same as above.

Solutions

6. Are $\{X_1, X_3\}$ and X_2 independent?
   Yes. Since they are multivariate normally distributed and the pairwise correlation between X_2 and $\{X_1, X_3\}$ is 0, they are independent.

7. Are $X_1 + X_2$ and $X_1 - X_2$ independent?
   No, the covariance between $X_1 + X_2$ and $X_1 - X_2$ is nonzero (see next question). Also, both terms involve X_1 and X_2 so there's no reason to expect them to be independent.

Solutions

8. Let $\mathbf{Y} = \mathbf{AX} + \mathbf{a}$ where $\mathbf{A} = \begin{pmatrix} 1 & 1 & 0 \\ 1 & -1 & 0 \end{pmatrix}$ and $\mathbf{a} = \begin{pmatrix} 0 \\ 1 \end{pmatrix}$.

   Find the distribution of $\mathbf{Y}$.

   $$\mathbf{Y} \sim N(\mathbf{A}\mu + \mathbf{a}, \mathbf{A}\mathbf{\Sigma}\mathbf{A}^T) = N\left(\begin{pmatrix} 1 \\ 0 \end{pmatrix}, \begin{pmatrix} 4 & 2 \\ 2 & 4 \end{pmatrix} \right)$$

9. What is the distribution of $W = (\mathbf{X} - \mu)(\mathbf{X} - \mu)^T$?
   A 3-dimensional Wishart distribution with one degree of freedom and matrix $\mathbf{\Sigma}$.

Reference Guide

- *Matrix Algebra from a Statistician's Perspective* - Harville
- *Matrix Algebra: Theory, Computations, and Applications in Statistics* - Gentle
- *The Matrix Cookbook* - Petersen and Pedersen
- *Applied Multivariate Statistical Analysis* - Johnson and Wichern
- *Modern Multivariate Statistical Techniques* - Izenman